

Deep Learning for Cybersecurity

Detecting botnets using ANNs

Barton Rhodes

Acknowledgements

This talk is based on a paper: “Predicting Domain Generation Algorithms with Long Short Term Memory Networks” by Woodbridge et al ([arXiv:1611.00791](https://arxiv.org/abs/1611.00791))

Predicting Domain Generation Algorithms with Long Short-Term Memory Networks

Jonathan Woodbridge, Hyrum S. Anderson, Anjum Ahuja, and Daniel Grant

{jwoodbridge, hyrum, aahuja, dgrant}@endgame.com

Endgame, Inc.

Arlington, VA 22201

Abstract—Various families of malware use domain generation algorithms (DGAs) to generate a large number of pseudo-random domain names to connect to a command and control (C2) server. In order to block DGA C2 traffic, security organizations must first discover the algorithm by reverse engineering malware samples, then generate a list of domains for a given seed. The domains are then either preregistered, sink-holed or published in a DNS blacklist. This process is not only tedious, but can be readily circumvented by malware authors. An alternative

server from which it can update, upload gathered intelligence, or pursue other malicious activities. The malicious actor only needs to register a small number of these domains to be successful. However, all the domains must be sinkholed, registered, or blacklisted before they go into use in order to preemptively defeat such an attack. This defense becomes increasingly difficult as the rate of dynamically generated domains increases.

Additional thanks to Miles Rufat-Latre, Mark Sliva, and Joewie Koh for working with me on the PyTorch implementation and to Jason Mancuso for being an awesome gent

Code

The code for the demo is available at:

<https://github.com/bmorphism/woodbridge-lstm> (you'll probably want a GPU!)

These slides:

<https://github.com/bmorphism/talks/tree/master/2017-11-14-deep-learning-cybersecurity>

You may also be interested in my previous KerasR talk for Denver R User Group:

<https://github.com/bmorphism/talks/tree/master/2017-08-10-keras>


Problem

- Various families of malware use domain generation algorithms (DGAs) to generate a large number of pseudo-random domain names deterministically
- Register a command and control (C2) server on a small portion of the domain names
- Have botnet try all possible domain names until the C2 is found
- To block DGA C2 traffic, security organizations must first discover the algorithm by reverse engineering malware samples, then generate a list of domains for a given seed
- Tedious and largely ineffective using traditional methods

Solution - supervised learning!

Given examples of domains generated by algorithms and common non-botnet domains, learn how to identify one from another.

Hold-out test set to evaluate model performance.



What would you say if I told you there is a app on the market that tell you if you have a hotdog or not a hotdog.

Previous solutions

- rely on observing statistical properties, such as distributions of bigrams
- in the case of classical ML, like Random Forest, rely on hand-crafted features
- require large windows or don't generalize as well

Examples:

- Hidden Markov Models
- Jaccard's similarity on n-grams
- Logistic Regression on bigrams
- Random Forest classifier

Why Neural Networks are better

Universal Approximation Theorem:

In the [mathematical](#) theory of [artificial neural networks](#), the **universal approximation theorem** states^[1] that a [feed-forward](#) network with a single hidden layer containing a finite number of [neurons](#) (i.e., a [multilayer perceptron](#)), can approximate [continuous functions](#) on [compact subsets](#) of \mathbf{R}^n , under mild assumptions on the activation function. (Source: [Wikipedia](#))

<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

Feature engineering


In particular, the manually crafted features of the random forest DGA classifier include the following:

- length of domain name,
- entropy of character distribution in domain name,
- vowel to consonant ratio,
- Alexa 1M n -gram frequency distribution co-occurrence count, where $n = 3, 4$ or 5 ,
- n -gram normality score, and
- *meaningful characters ratio*.

What is a neural network?

<http://deeplearning.ai>

<http://fast.ai>



What is keras?

<https://keras.io/>


Core Layers

Dense
Activation
Dropout
Flatten
Reshape
Permute
RepeatVector
Lambda
ActivityRegularization
Masking

Recurrent Layers

RNN
SimpleRNN
GRU
LSTM
ConvLSTM2D
SimpleRNNCell
GRUCell
LSTMCell
StackedRNNCells
CuDNNGRU
CuDNNLSTM

Simple example



```
from keras.models import Sequential
```

```
from keras.layers import Dense
```

```
model = Sequential()
```

```
model.add(Dense(512, activation='relu',  
input_shape=(dimData,)))
```


```
model.add(Dense(512, activation='relu'))
```

```
model.add(Dense(nClasses,  
activation='softmax'))
```


<https://keras.io/#getting-started-30-seconds-to-keras>

A word about activation functions


Sigmoid



TanH



ReLU




<https://medium.com/the-theory-of-everything/understanding-activation-functions-in-neural-networks-9491262884e0>

Data

- Alexa Top 1M domains
- OSINT Domain-Generation Algorithms feed (as well as domain generation snippets published with the paper's code)

For the DGA (malicious) domains, 750k examples across 30 botnet classes.

Architecture



Embedding

representation of the ***semantics*** of a word, efficiently encoding semantic information that might be relevant to the task at hand

$$q_{\text{mathematician}} = \left[\begin{array}{c} \text{can run} \quad \text{likes coffee} \quad \text{majored in Physics} \\ \underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad} \\ 2.3 \quad , \quad 9.4 \quad , \quad -5.5 \quad , \dots \end{array} \right]$$

$$q_{\text{physicist}} = \left[\begin{array}{c} \text{can run} \quad \text{likes coffee} \quad \text{majored in Physics} \\ \underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad} \quad \underbrace{\quad\quad\quad} \\ 2.5 \quad , \quad 9.1 \quad , \quad 6.4 \quad , \dots \end{array} \right]$$

Visualizing features is difficult






Fig. 8: Two-dimensional linear projection (PCA) of the embedded character vectors learned by the LSTM binary classifier. Note that the model groups characters by similar effect on the LSTM layer's states and the subsequent model loss.




LSTM




Legend

- unweighted connection
- weighted connection
- - - connection with time-lag
- branching point
- ⊙ multiplication
- ⊕ sum over all inputs
- σ gate activation function (always sigmoid)
- g input activation function (usually tanh)
- h output activation function (usually tanh)

Dropout - prevent overfitting



(a) Standard Neural Net



(b) After applying dropout.

Keras

```
1 from keras.preprocessing import pad_sequences
2 from keras.models import Sequential
3 from keras.layers.core import Dense
4 from keras.layers.core import Dropout
5 from keras.layers.core import Activation
6 from keras.layers.embeddings import Embedding
7 from keras.layers.recurrent import LSTM
8
9 model=Sequential()
10 model.add(Embedding(max_features,
11                     128,
12                     input_length=75))
13 model.add(LSTM(128))
14 model.add(Dropout(0.5))
15 model.add(Dense(1))
16 model.add(Activation('sigmoid'))
17
18 model.compile(loss='binary_crossentropy',
19              optimizer='rmsprop')
20
21 # Pad sequence where sequences are case
22 # insensitive characters encoded to
23 # integers from 0 to number of valid
24 # characters
25 X_train=sequence.pad_sequences(X_train,
26                               maxlen=75)
27
28 # Train where y_train is 0-1
29 model.fit(X_train, y_train,
30         batch_size=batch_size, nb_epoch=1)
```

Fig. 2: Binary LSTM Code

```
1 from keras.preprocessing import pad_sequences
2 from keras.models import Sequential
3 from keras.layers.core import Dense
4 from keras.layers.core import Dropout
5 from keras.layers.core import Activation
6 from keras.layers.embeddings import Embedding
7 from keras.layers.recurrent import LSTM
8
9 model=Sequential()
10 model.add(Embedding(max_features,
11                     128,
12                     input_length=75))
13 model.add(LSTM(128))
14 model.add(Dropout(0.5))
15 # nb_classes is the number of classes in
16 # the training set
17 model.add(Dense(nb_classes))
18 model.add(Activation('softmax'))
19
20 model.compile(loss='categorical_crossentropy',
21              optimizer='rmsprop')
22
23 # Pad sequence where sequences are case
24 # insensitive characters encoded to
25 # integers from 0 to number of valid
26 # characters
27 X_train=sequence.pad_sequences(X_train,
28                               maxlen=75)
29
30 # Train where y_train is one-hot encoded for
31 # each class
32 model.fit(X_train, y_train,
33         batch_size=batch_size, nb_epoch=1)
```

Fig. 3: Multiclass LSTM Code

Metrics

$$\text{Precision} = \frac{\sum \text{True Positive}}{\sum \text{True Positive} + \sum \text{False Positive}}$$

$$\text{Recall} = \frac{\sum \text{True Positive}}{\sum \text{True Positive} + \sum \text{False Negative}}$$

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

$$\text{TPR} = \frac{\sum \text{True Positive}}{\sum \text{True Positive} + \sum \text{False Negative}}$$

$$\text{FPR} = \frac{\sum \text{False Positive}}{\sum \text{False Positive} + \sum \text{True Negative}}$$

Results




TABLE II: Precision, Recall and F_1 Score for Binary Classifiers

Domain Type	Precision				Recall				F_1 Score				Support
	HMM	Features	Bigram	LSTM	HMM	Features	Bigram	LSTM	HMM	Features	Bigram	LSTM	
Alexa	0.8300	0.9400	0.9700	0.9900	1.0000	1.0000	1.0000	1.0000	0.9100	0.9700	0.9900	0.9900	300064
Cryptolocker	1.0000	1.0000	1.0000	1.0000	0.9000	0.9800	0.9700	0.9900	0.9500	0.9900	0.9900	0.9900	1799
P2P Gameover Zeus	1.0000	1.0000	1.0000	1.0000	0.9900	1.0000	1.0000	1.0000	0.9900	1.0000	1.0000	1.0000	298
Post Tovar GOZ	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	19863
Volatile Cedar / Explosive	0.0000	1.0000	1.0000	1.0000	0.0000	0.4600	0.4900	0.9900	0.0000	0.6300	0.6600	1.0000	294
banjori	1.0000	1.0000	1.0000	1.0000	0.5900	0.9400	1.0000	1.0000	0.7400	0.9700	1.0000	1.0000	121678
bedep	1.0000	1.0000	1.0000	1.0000	0.8100	1.0000	1.0000	1.0000	0.8900	1.0000	1.0000	1.0000	53
beebone	0.0000	1.0000	1.0000	1.0000	0.0000	1.0000	0.9700	1.0000	0.0000	1.0000	0.9900	1.0000	65
corebot	1.0000	1.0000	1.0000	1.0000	0.5900	1.0000	1.0000	0.9600	0.7400	1.0000	1.0000	0.9800	81
cryptowall	1.0000	1.0000	1.0000	1.0000	0.1100	0.0600	0.1400	0.1200	0.1900	0.1100	0.2500	0.2100	29
dircrypt	1.0000	1.0000	1.0000	1.0000	0.9100	0.9200	0.9600	0.9600	0.9500	0.9600	0.9800	0.9800	150
dyre	1.0000	1.0000	1.0000	1.0000	1.0000	1.0000	0.9900	1.0000	1.0000	1.0000	0.9900	1.0000	2389
fobber	1.0000	1.0000	1.0000	1.0000	0.8900	0.9600	0.9700	0.9700	0.9400	0.9800	0.9800	0.9900	181
geodo	1.0000	1.0000	1.0000	1.0000	0.9100	1.0000	0.9900	0.9900	0.9500	1.0000	1.0000	1.0000	173
hesperbot	1.0000	1.0000	1.0000	1.0000	0.8300	0.7700	0.8500	0.9700	0.9100	0.8700	0.9200	0.9800	58
matsnu	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000	14
murofet	1.0000	1.0000	1.0000	1.0000	0.9200	1.0000	0.9900	1.0000	0.9600	1.0000	1.0000	1.0000	4292
necurs	1.0000	1.0000	1.0000	1.0000	0.8800	0.8400	0.9400	0.9600	0.9400	0.9100	0.9700	0.9800	1232
nymaim	1.0000	1.0000	1.0000	1.0000	0.8000	0.5600	0.7300	0.8000	0.8900	0.7200	0.8500	0.8900	1815
pushdo	1.0000	1.0000	1.0000	1.0000	0.6600	0.4700	0.5600	0.6000	0.7900	0.6400	0.7200	0.7500	507
pykspa	1.0000	1.0000	1.0000	1.0000	0.7200	0.5400	0.7700	0.9000	0.8400	0.7000	0.8700	0.9500	4250
qakbot	1.0000	1.0000	1.0000	1.0000	0.9100	0.9600	0.9600	0.9800	0.9500	0.9800	0.9800	0.9900	1517
ramnit	1.0000	1.0000	1.0000	1.0000	0.8800	0.9100	0.9400	0.9600	0.9400	0.9500	0.9700	0.9800	27439
ranbyus	1.0000	1.0000	1.0000	1.0000	0.9000	1.0000	0.9800	0.9800	0.9500	1.0000	0.9900	0.9900	2625
shifu	1.0000	1.0000	1.0000	1.0000	0.7200	0.2100	0.6600	0.7700	0.8400	0.3500	0.8000	0.8700	697
shiotob/urlzone/bebloh	1.0000	1.0000	1.0000	1.0000	0.9000	0.9700	0.9500	0.9800	0.9500	0.9900	0.9700	0.9900	3031
simda	1.0000	1.0000	1.0000	1.0000	0.5600	0.0800	0.4000	0.9200	0.7100	0.1400	0.5800	0.9600	4449
suppobox	1.0000	0.0000	1.0000	1.0000	0.0100	0.0000	0.0000	0.3200	0.0200	0.0000	0.0100	0.4800	298
symmi	0.0000	1.0000	1.0000	1.0000	0.0000	1.0000	0.7900	0.6900	0.0000	1.0000	0.8800	0.8200	18
tempedreve	1.0000	1.0000	1.0000	1.0000	0.7600	0.5700	0.8500	0.7700	0.8600	0.7300	0.9200	0.8700	74
tinba	1.0000	1.0000	1.0000	1.0000	0.8900	0.9800	0.9700	0.9900	0.9400	0.9900	0.9900	0.9900	18505
Micro Average	0.9008	0.9647	0.9826	0.9942	0.8815	0.9639	0.9848	0.9937	0.8739	0.9593	0.9851	0.9906	16708
Macro Average	0.8655	0.9335	0.9668	0.9674	0.6787	0.7477	0.8006	0.8571	0.7335	0.7929	0.8468	0.8913	16708

TABLE IV: Precision, Recall and F_1 Score for Multiclass Classifiers

Domain Type	Precision			Recall			F_1 Score			Support
	Features	Bigram	LSTM	Features	Bigram	LSTM	Features	Bigram	LSTM	
Alexa	0.914	0.980	0.990	0.960	0.990	1.000	0.940	0.988	0.990	199978
Cryptolocker	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	1189
P2P Gameover Zeus	0.000	0.343	0.327	0.000	0.288	0.217	0.000	0.308	0.247	196
Post Tovar GOZ	0.941	1.000	1.000	1.000	1.000	1.000	0.970	1.000	1.000	13185
Volatile Cedar / Explosive	0.000	1.000	0.987	0.000	1.000	0.980	0.000	1.000	0.980	200
banjori	0.900	0.990	1.000	0.938	1.000	1.000	0.920	1.000	1.000	81281
bedep	0.000	0.000	0.943	0.000	0.000	0.107	0.000	0.000	0.187	34
beebone	1.000	1.000	1.000	0.560	1.000	1.000	0.713	1.000	1.000	42
corebot	0.000	1.000	1.000	0.000	0.980	0.990	0.000	0.990	0.993	54
cryptowall	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	15
dircrypt	0.000	0.083	0.000	0.000	0.010	0.000	0.000	0.020	0.000	100
dyre	0.985	0.988	1.000	1.000	0.988	1.000	0.991	0.988	1.000	1600
fobber	0.000	0.000	0.177	0.000	0.000	0.023	0.000	0.000	0.040	121
geodo	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	114
hesperbot	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	36
matsnu	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	9
murofet	0.883	0.643	0.783	0.066	0.542	0.700	0.122	0.590	0.737	2845
necurs	0.000	0.000	0.643	0.000	0.000	0.093	0.000	0.000	0.160	827
nymaim	0.000	0.390	0.477	0.000	0.113	0.190	0.000	0.175	0.267	1222
pushdo	0.000	0.770	0.853	0.000	0.588	0.640	0.000	0.665	0.730	339
pykspa	0.000	0.788	0.910	0.000	0.593	0.713	0.000	0.675	0.800	2827
qakbot	0.000	0.590	0.590	0.000	0.232	0.387	0.000	0.338	0.463	993
ramnit	0.566	0.637	0.770	0.654	0.763	0.850	0.605	0.690	0.810	18308
ranbyus	0.439	0.000	0.450	0.000	0.000	0.517	0.001	0.000	0.460	1736
shifu	0.000	0.037	0.560	0.000	0.003	0.570	0.000	0.007	0.553	465
shiotob/urlzone/bebloh	0.000	0.965	0.973	0.000	0.853	0.907	0.000	0.907	0.940	2016
simda	0.000	0.840	0.930	0.000	0.750	0.977	0.000	0.792	0.950	2955
suppobox	0.000	0.392	0.833	0.000	0.062	0.517	0.000	0.112	0.627	197
symmi	0.000	0.625	0.913	0.000	0.117	0.857	0.000	0.200	0.883	11
tempedreve	0.000	0.043	0.000	0.000	0.010	0.000	0.000	0.018	0.000	50
tinba	0.821	0.735	0.910	0.923	0.802	0.990	0.869	0.767	0.950	12332
Micro Average	0.851	0.933	0.963	0.888	0.944	0.970	0.867	0.940	0.963	11138
Macro Average	0.240	0.479	0.614	0.197	0.409	0.523	0.198	0.427	0.541	11138

Demo

Putting it all together in Keras

Conclusions

- by automating feature extraction and working with better representations in neural networks, one can accomplish state-of-the-art results on long-standing cybersecurity tasks
- Keras is a great library to quickly and easily reproduce a lot of papers (except ones you'll need PyTorch for
<https://medium.com/intuitionmachine/pytorch-dynamic-computational-graphs-and-modular-deep-learning-7e7f89f18d1>)
- can be productionised to neutralize botnets, including new kinds, before they spread

Denver DL Study Group

Small meeting @ 2 pm at [Denver Bicycle Café](#) every Sunday.

Currently working through:

- Andrew Ng's Deep Learning Specialization on Coursera
- fast.ai Deep Learning for Coders (Part II)
- Passenger screening algorithm Kaggle competition

Contact me at [@bmorphism](#) on Twitter or over email at b@bmorphism.us to be added to our Slack.

Thanks!